

+ 24

# LLVM's Realtime Safety Revolution:

Tools for Modern Mission Critical Systems

DAVID TREVELYAN &  
CHRISTOPHER APPLE



20  
24



# Chris Apple

- 10-year veteran of the audio industry
- Previously Dolby, Roblox, Spatial Inc.
- Currently: layabout



# David Trevelyan

- 15 years' experience in academia, startups and big tech
- Previously TikTok, Imperial College London
- Currently: audio software consultant



**Authors of RealtimeSanitizer**

# Special thanks



**Alistair Barker**

Co-author of RealtimeSanitizer



**Doug Wyatt**

Inventor/author of performance constraints

# AGENDA

1. Real-time programming
2. Existing strategies
3. RealtimeSanitizer
4. Performance constraints
5. Comparing and contrasting

50%

50%

# AGENDA

1. **Real-time programming**
2. Existing strategies
3. RealtimeSanitizer
4. Performance constraints
5. Comparing and contrasting

Real-time programs must guarantee response within specified time constraints, often referred to as “deadlines”.

1. Provide the right answer...

2. ... in the right time.



**Missed deadlines**

Real-time programming





**POP!**









# Real-time code in C++

Real-time programming

```
void process (float * audio)
{
    // You have 11.6 ms from now, GO!!!!
    ...
}
```

```
void dispatch (Data& audio)
{
    // Fly a rocket...
    ...
}
```

# Worst case execution time must be...

1. deterministic,
2. known in advance,
3. independent of input data, and
4. shorter than the given deadline [1].

*[1] Timur Doumler - Real Time Programming with the C++ Standard Library*



# How long is a malloc?

10 ns

1,000,000 ns

90%

< 1%

How long is a real-time callback?

Milliseconds

Microseconds

Nanoseconds

# Nondeterministic execution time

1. System calls
2. Allocations
3. Mutex locks/unlocks
4. Thrown exceptions
5. Indefinite waits (e.g. loops, infinite loops)
6. ...

**[PROHIBITED]**

nonblocking



**What do violations  
look like?**

Real-time programming

Sometimes it's  
*obvious*

```
void dispatch()  
{  
    mutex_.lock();  
    ...  
}
```



```
void dispatch()  
{  
    x = make_unique<Foo>();  
    ...  
}
```





**Sometimes it's**  
hidden

```
void process_audio()  
{  
    numbers_[1] = 2;  
    ...  
}
```



```
void dispatch()
{
    auto const x = input_array();
    auto const y = output_array();

    post_report([x, y](auto & data) {
        data.input = x;
        data.output = y;
    });
}
```



```
void process_audio()  
{  
    fftw_execute(plan);  
    ...  
}
```



How can we be  
confident that  
**our** code is  
real-time safe?

# AGENDA

1. Real-time programming
- 2. Existing strategies**
3. RealtimeSanitizer
4. Performance constraints
5. Comparing and contrasting

# Existing strategies

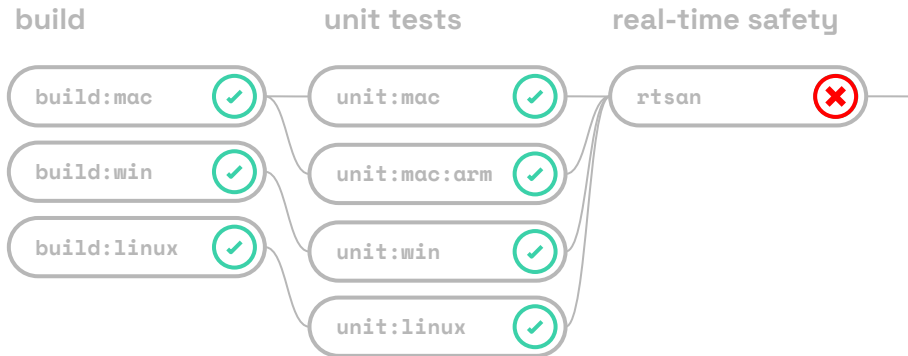
- Shared experience
  - Code review
  - Profilers and debuggers
  - `static_assert`
  - Documentation
- Getting experience takes a **long time**
  - Code review is prone to **human error**
  - Profiling/debugging is a **manual process**
  - Static assertions are **limited**
  - Documentation **goes out of date**
  - *What about **pre-built dependencies**?*

**What if we had a tool  
that could  
simply tell us?**



# A nice tool would...

- Assess real-time safety
- Detect a wide range of violations
- ...even from third-party and pre-compiled dependencies
- Point to any problematic code
- Be able to fail a CI pipeline



# LLVM 20



1. RealtimeSanitizer
2. Performance constraints

# AGENDA

1. Real-time programming
2. Existing strategies
- 3. RealtimeSanitizer**
4. Performance constraints
5. Comparing and contrasting

# Using sanitizers

```
#include <vector>

int main()
{
    auto v = std::vector<int> (16);
    return v[16];
}
```

```
> clang -fsanitize=address main.cpp
```

```
> ./a.out
```

```
==98481==ERROR: AddressSanitizer: heap-buffer-overflow on address 0x000105701320 at pc 0x000102d8
```

```
READ of size 4 at 0x000105701320 thread T0
```

```
#0 0x5770e099d6c0 in main /app/example.cpp:6:12
```

```
#1 0x7dbed6c29d8f (/lib/x86_64-linux-gnu/libc.so.6+0x29d8f)
```

```
#2 0x7dbed6c29e3f in __libc_start_main (/lib/x86_64-linux-gnu/libc.so.6+0x29e3f)
```

```
#3 0x5770e08ba394 in _start (/app/output.s+0x2c394)
```

```
0x606000000300 is located 0 bytes after 64-byte region [0x6060000002c0,0x606000000300)
```

```
allocated by thread T0 here:
```

```
#0 0x104c7f954 in _Znwm asan_new_delete.cpp:120
```

```
#1 0x1044294cc in void* std::__1::__libcxx_operator_new[abi:v160006]<unsigned long>(unsigned
```

```
#2 0x104429450 in std::__1::__libcxx_allocate[abi:v160006](unsigned long, unsigned long)+0x44
```

# Using RealtimeSanitizer (RTSan)

```
float process(float x) [[clang::nonblocking]]
{
    auto const y = std::vector<float> (16);
    ...
}
```

```
> clang++ -fsanitize=realtime main.cpp
```

```
> ./a.out
```

```
==86660==ERROR: RealtimeSanitizer: unsafe-library-call
```

```
Intercepted call to real-time unsafe function `malloc` in real-time context!
```

```
#0 0x103184cfc in malloc rtsan_interceptors.cpp:225
```

```
#1 0x18cb16524 in operator new(unsigned long)+0x1c
```

```
...
```

```
#10 0x102c02b8c in std::__1::vector<float, std::__1::allocator<float>>::vector
```

```
#11 0x102c02b38 in process(float)+0x28 /app/example.cpp:6:14
```

```
#12 0x102c02c00 in main+0x1c /app/example.cpp:12:5
```

# Two steps

1. Attribute real-time functions with `[[clang::nonblocking]]`
2. Compile and link with `-fsanitize=realtime`





# How RealtimeSanitizer works

RealtimeSanitizer

# Tracking and interception

lightweight runtime library

```
void __rtsan_realtime_enter() { ... }  
void __rtsan_realtime_exit() { ... }
```

```
INTERCEPTOR (void *, malloc, size_t size)  
{  
    if (is_in_realtime_context()):  
        print_stack_and_die("malloc");  
  
    return REAL(malloc)(size);  
}
```

LLVM Intermediate  
Representation (IR)

# Real-time context signaling

compilation step

```
int dispatch() [[clang::nonblocking]] {  
    return calculate_result();  
}
```



-fsanitize=realtime

```
define noundef i32 @_Z8dispatchv() #1 {  
    call void @__rtsan_realtime_enter()  
    %1 = call noundef i32 @_Z16calculate_resultv()  
    call void @__rtsan_realtime_exit()  
    ret i32 %1  
}
```



# Interception

Interception allows us to:

- detect any function call
- assert not in real-time context
- error or behave normally

**RealtimeSanitizer aims to intercept anything that could block**

```
INTERCEPTOR (void *, malloc, size_t size) {  
    __rtsan_expect_not_realtime("malloc");  
    return REAL(malloc)(size);  
}
```

```
INTERCEPTOR (void, free, void * ptr) {  
    __rtsan_expect_not_realtime("free");  
    return REAL(free)(ptr);  
}
```

```
INTERCEPTOR (int, pthread_mutex_lock, pthread_mutex_t * mutex) {  
    __rtsan_expect_not_realtime("pthread_mutex_lock");  
    return REAL(pthread_mutex_lock)(mutex);  
}
```

```
INTERCEPTOR (int, pthread_mutex_unlock, pthread_mutex_t * mutex) {  
    __rtsan_expect_not_realtime("pthread_mutex_unlock");  
    return REAL(pthread_mutex_unlock)(mutex);  
}
```

```
INTERCEPTOR (int, pthread_cond_signal, pthread_cond_t * cond) {  
    __rtsan_expect_not_realtime("pthread_cond_signal");  
    return REAL(pthread_cond_signal)(cond);  
}
```

## Allocations

User space

User C++ Application

```
auto w = std::make_unique<Widget>();  
throw std::runtime_error(msg);  
auto v = std::vector<float>(4);  
return std::string(err);
```


C++ Runtime Library (**libc++**)

```
operator new  
__cxa_allocate_exception
```

C Standard Library (**libc**)

**malloc**


RealtimeSanitizer  
interceptors



Kernel System Call Interface (SCI)

Kernel space

Kernel Internals



## Memory allocation

`malloc, calloc`

`realloc, reallocf,`

`valloc, aligned_alloc`

`free`

`posix_memalign`

## Threads & sleep

`pthread_create, pthread_join`

`pthread_mutex_lock, pthread_mutex_unlock`

`pthread_cond_signal, pthread_cond_broadcast`

`OSSpinLockLock, os_unfair_lock_lock`

`sleep, usleep, nanosleep`

## Filesystem & streams

`open, openat, creat, close,`

`fopen, fopenat, fclose, fread, fwrite,`

`puts, fputs, read, write, writev, readv,`

`pwrite, pread`

## Sockets

`socket`

`send, sendto, sendmsg`

`recv, recvfrom, recvmsg`

`shutdown`

```
float process(float x)  [[clang::nonblocking]]
{
    auto y = std::vector<float> (16);
    ...
}
```

```
> clang++ -fsanitize=realtime main.cpp
```

```
> ./a.out
```

```
==86660==ERROR: RealtimeSanitizer: unsafe-library-call
```

```
Intercepted call to real-time unsafe function `malloc` in real-time context!
```

```
#0 0x103184cfc in malloc rtsan_interceptors.cpp:225
```

```
#1 0x18cb16524 in operator new(unsigned long)+0x1c
```

```
...
```

```
#10 0x102c02b8c in std::__1::vector<float, std::__1::allocator<float>>
```

```
#11 0x102c02b38 in process(float)+0x28 /app/example.cpp:6:14
```

```
#12 0x102c02c00 in main+0x1c /app/example.cpp:12:5
```

# AGENDA

1. Real-time programming
2. Existing strategies
3. RealtimeSanitizer
- 4. Performance constraints**
5. Comparing and contrasting

# Performance constraint compilation flags

-Wfunction-effects

-Wperf-constraint-implies-noexcept

# Performance constraint function attributes

`[[clang::nonallocating]]`

`[[clang::nonblocking]]`

## [[clang::nonallocating]]

1. Add `[[clang::nonallocating]]` to a function.
2. Compile with warning flags  
`-Wfunction-effects`  
`-Wperf-constraint-implies-noexcept`

```
void process()  
[[clang::nonallocating]] {  
    float* v = new float[100];  
}
```

warning: '**nonallocating**' function **must not allocate or deallocate memory**

`[-Wfunction-effects]`

```
16 |    float* v = new float[100];
```



## [[clang::nonblocking]]

1. Add `[[clang::nonblocking]]` to a function.
2. Compile with warning flags

`-Wfunction-effects`

`-Wperf-constraint-implies-noexcept`

```
void dispatch()  
[[clang::nonblocking]] {  
    m.lock();  
}
```

```
warning: 'nonblocking' function must not call  
non-'nonblocking' function 'std::mutex::lock'  
[-Wfunction-effects]  
5 |     m.lock();
```

# Performance constraints attributes

`[[clang::nonblocking]]`

implies

`[[clang::nonallocating]]`

requires

`noexcept`



```
void dispatch() noexcept  
[[clang::nonblocking]] {  
    float* ptr = new float;  
}
```

```
warning: 'nonblocking' function must  
not allocate or deallocate memory  
[-Wfunction-effects]  
4 | float* ptr = new float;
```

[[clang::nonblocking]]

implies

[[clang::nonallocating]]

requires

noexcept

```
void foo()  
[[clang::nonallocating]];
```

warning: '**nonallocating**'  
function should be declared  
**noexcept**

```
3 | void foo()
```

-Wperf-constraint-implies-noexcept

[[clang::nonblocking]]

implies



[[clang::nonallocating]]

requires



noexcept

# Performance constraints attributes

`[[clang::nonblocking]]`

implies

`[[clang::nonallocating]]`

requires

`noexcept`





# Function restrictions

Performance Constraints

Normal functions  
can call...

`[[nonblocking]]`

`[[nonallocating]]`

no attributes

`[[nonallocating]]`  
can call...

`[[nonblocking]]`

`[[nonallocating]]`

~~no attributes~~

`[[nonblocking]]`  
can call...

`[[nonblocking]]`

~~`[[nonallocating]]`~~

~~no attributes~~

# Function call restrictions

```
void unknown_function() [[clang::nonblocking]]; Missing[[nonblocking]]
```

```
void process() noexcept [[clang::nonblocking]] {  
    unknown_function();  
}
```



warning: 'nonblocking' function must not call non-'nonblocking' function 'unknown\_function' [-Wfunction-effects]

[[nonblocking]] can ONLY call other functions marked [[nonblocking]]



**Functions may only  
call their constraint,  
or stricter**

`[[clang::nonblocking]]`

implies

`[[clang::nonallocating]]`

## Does it compile?

```
void foo() [[clang::nonblocking]];
```

```
void process() [[clang::nonblocking]]  
{  
    foo();  
}
```



[[clang::nonblocking]]  
↓ implies  
[[clang::nonallocating]]

## Does it compile?

```
void foo();
```

```
void process() [[clang::nonblocking]]  
{  
    foo();  
}
```



```
main.cpp:6:3: warning: 'nonblocking'  
function must not call non-'nonblocking'  
function 'foo' [-Wfunction-effects]  
    6 |    foo();
```

[[clang::nonblocking]]  
↓ implies  
[[clang::nonallocating]]

## Does it compile?

```
void foo() [[clang::nonallocating]];  
  
void process() [[clang::nonblocking]]  
{  
    foo();  
}
```



```
main.cpp:6:3: warning: 'nonblocking' function  
must not call non-'nonblocking' function  
'foo' [-Wfunction-effects]  
    6 |   foo();
```

`[[clang::nonblocking]]`

implies

`[[clang::nonallocating]]`

## Does it compile?

```
void foo() [[clang::nonblocking]];
```

```
void process() [[clang::nonallocating]]  
{  
    foo();  
}
```





# Constraint inference

Performance Constraints

## Inference - Same TU

```
int defined_here() {}
```

```
void process() noexcept  
[[clang::nonblocking]] {  
    defined_here();  
}
```





```
int defined_here() {  
    mutex.lock();  
}
```

```
void process() noexcept  
[[clang::nonblocking]] {  
    defined_here();  
}
```

## Inference - Same TU

```
main.cpp:9:3: warning:  
'nonblocking' function must not  
call non-'nonblocking' function  
'defined_here' [-Wfunction-effects]
```

```
9 |     defined_here();
```

```
main.cpp:17:3: note: function  
cannot be inferred 'nonblocking'  
because it calls non-'nonblocking'  
function 'std::mutex::lock'
```

```
17 |     mutex.lock();
```





```
// third_party.h
```

```
void defined_elsewhere();
```

```
// main.cpp
```

```
void process() noexcept  
[[clang::nonblocking]] {  
    defined_elsewhere();  
}
```

## Inference - Different TU

```
main.cpp:5:3: warning: 'nonblocking'  
function must not call non-'nonblocking'  
function 'defined_elsewhere'  
[-Wfunction-effects]
```

```
5 |     defined_elsewhere();
```

```
main.cpp:2:5: note: declaration cannot be  
inferred 'nonblocking' because it has no  
definition in this translation unit
```

```
2 | int defined_elsewhere();
```

# Re-declaration

```
// third_party.h
```

```
void defined_elsewhere();
```

```
// main.cpp
```

```
void defined_elsewhere() [[clang::nonblocking]];
```

```
void process() [[clang::nonblocking]] {
```

```
    defined_elsewhere();
```

```
}
```



## Re-declaration

```
// third_party.h
```

```
void defined_elsewhere();
```

```
// main.cpp
```

```
void defined_elsewhere() [[clang::nonblocking]];
```

```
void process() [[clang::nonblocking]] {
```

```
    defined_elsewhere();
```

```
}
```

```
// elsewhere.cpp
```

```
void defined_elsewhere()
```

```
{
```

```
    mutex.lock();
```

```
    ...
```

```
}
```



**This will compile without warnings, even though it is incorrect!**

# Summary of performance constraints

```
> clang++ -Wfunction-effects -Wperf-constraint-implies-noexcept main.cpp
```

```
void process() noexcept [[clang::nonblocking]] {
```

```
    float* f = new float;
```

```
    // error: 'nonblocking' function must not allocate or deallocate memory
```

```
    auto g = foo();
```

```
    // error: 'nonblocking' function must not call non-'nonblocking' function 'foo'
```

```
    static int x = 0;
```

```
    // error: 'nonblocking' function must not have static locals
```

```
}
```

# AGENDA

1. Real-time programming
2. Existing strategies
3. RealtimeSanitizer
4. Performance constraints
5. **Comparing and contrasting**

# INTERLUDE

**WARNING!**

**Neither RTSan nor the  
perf. constraints  
attributes can fully  
guarantee real-time  
safety**



# RTSan blind spots

- No guarantee of processor time.
- No guarantee your code runs faster than allotted time.
- No detection of hand-written assembly system calls.
- Not all libc wrapper functions implemented.
- No detection of nondeterministic loops.
  - Infinite loops
  - Nondeterministic loops (CAS)



# RTSan blind spots

- No guarantee of processor time.
- No guarantee your code runs faster than allotted time.
- No detection of hand-written assembly system calls.
- Not all libc wrapper functions implemented.
- No detection of nondeterministic loops.
  - Infinite loops
  - Nondeterministic loops (CAS)



# Perf. constraints blind spots

- No guarantee of processor time.
- No guarantee your code runs faster than allotted time.
- Misdeclared functions.



# Perf. constraints blind spots

- No guarantee of processor time.
- No guarantee your code runs faster than allotted time.
- Misdeclared functions.



Neither RTSan nor the  
perf. constraints  
attributes can fully  
guarantee real-time  
safety

These tools  
make writing  
real-time code  
safer

# INTERLUDE OVER

# AGENDA

1. Real-time programming
2. Existing strategies
3. RealtimeSanitizer
4. Performance constraints
5. **Comparing and contrasting**



RealtimeSanitizer

Performance constraints

**Both**

**Detect real-time safety issues**

**Use `[[clang::nonblocking]]`**

**No real-time safety guarantee**

1. Run time vs compile time
2. False negatives and false positives
3. Cost
4. Disabling each tool
5. Using each tool outside of LLVM

# CONTRAST

1. **Run time vs compile time**
2. False negatives and false positives
3. Cost
4. Disabling each tool
5. Using each tool outside of LLVM

**CONTRAST**

RealtimeSanitizer

Run time

Performance constraints

Compile time

# RealtimeSanitizer

If your code runs, it is compliant.\*

**\* As long as you hit every path in your code.**

```
if (something_rare())  
    return process_that_allocates();
```

Recommendation:

1. Extensive unit testing with RealtimeSanitizer
2. QA testing with RealtimeSanitizer

# Performance constraints

If your code compiles without warnings, it is compliant.\*

**\* As long as you:**

1. **Pay attention to the warnings**
2. **Declare everything correctly, especially in third party dependencies**

Recommendation:

1. Compile with **-Werror=function-effects** and **-Werror=perf-constraint-implies-noexcept**
2. “Audit” 3rd party code with RealtimeSanitizer

# RealtimeSanitizer

**With RealtimeSanitizer enabled:**

- Extensive unit testing
- Extensive QA testing

# Performance constraints

**Compile with -Werror**

**“Audit” 3rd party code with  
RealtimeSanitizer**

1. **Run time vs compile time**
2. False negatives and false positives
3. Cost
4. Disabling each tool
5. Using each tool outside of LLVM

# CONTRAST

1. Run time vs compile time
2. **False negatives and false positives**
3. Cost
4. Disabling each tool
5. Using each tool outside of LLVM

# CONTRAST



**RealtimeSanitizer**

**Possible false negatives**

**Performance constraints**

**Possible false positives**

## Perf. constraints

```
main.cpp:8:3: error: 'nonblocking'  
function must not call non-'nonblocking'  
function 'std::vector<int>::push_back'
```

```
8 |    v.push_back(3);
```

push\_back real-time safe?

```
void dispatch(vector<int>& v)  
noexcept [[clang::nonblocking]]  
{  
    v.clear();           ...  
    v.push_back(3);      vector<int> v;  
                          v.reserve(512);  
                          dispatch(v);  
                          ...  
}
```

## RealtimeSanitizer



push\_back real-time safe?

```
void dispatch(vector<int>& v)
noexcept [[clang::nonblocking]]
{
    v.clear();           ...
    v.push_back(3);      vector<int> v;
}                        v.reserve(512);
                        dispatch(v);
                        ...
```

# Pre-reserved push\_back

## RTSan

- Did not error
- “This method is compliant on all the paths we hit!”

**Especially useful for probing third-party  
libraries and the STL!**

## Performance constraints

- False positive, showed an error when the code did not allocate.
- “This method is not always free of non-blocking calls!”

**Perf. constraints  
analysis has no  
false negatives**

**But it is prone to  
false positives**

RealtimeSanitizer  
has no false  
positives

But it is prone to  
false negatives

1. Run time vs compile time
2. **False negatives and false positives**
3. Cost
4. Disabling each tool
5. Using each tool outside of LLVM

# CONTRAST

1. Run time vs compile time
2. False negatives and false positives
- 3. Cost**
4. Disabling each tool
5. Using each tool outside of LLVM

**CONTRAST**



**RealtimeSanitizer**

**Run time cost**

**Performance constraints**

**Code conversion cost**



# RealtimeSanitizer run-time cost

Comparing and contrasting

# How does RTSan change your code?

```
int dispatch() [[clang::nonblocking]] {  
    __rtsan_realtime_enter();  
    ...fly a rocket code here...  
    __rtsan_realtime_exit();  
}
```

# How does RTSan change your code?

```
int dispatch() [[clang::nonblocking]] {  
    realtime_depth++; // thread-local integer  
    ...fly a rocket code here...  
    realtime_depth--;  
}
```

## [[nonblocking]]

- Every [[nonblocking]] function has an additional increment and decrement of an integer
- May interfere with inlining optimizations

## Non-constrained?

## How does RTSan change your code?

```
int ui_thread() {
```

```
    auto w = std::make_unique<Widget>(...);
```

```
}
```

# How does RTSan change your code?

```
int ui_thread() {
```

```
    Widget* w = (Widget*)malloc(sizeof(Widget));
```

```
}
```

# How does RTSan change your code?

```
int ui_thread() {  
  
    Widget* w = (Widget*)rtsan_malloc(sizeof(Widget));  
  
}
```



# How does RTSan change your code?

```
int ui_thread() {  
    if (not __rtsan_in_dlsym())  
        if (__rtsan_initialized())  
            if (not __rtsan_in_realtime_context())  
                Widget* w = (Widget*)real_malloc(sizeof(Widget));  
}
```

## [[nonblocking]]

- Every [[nonblocking]] function has an additional increment and decrement of an integer
- May interfere with inlining optimizations

## Non-constrained

- A few additional function calls inserted between the intercepted call, and the real method being invoked
  - These functions are relatively light, checking state stored in integers

## [[nonblocking]]

- Every [[nonblocking]] function has an additional increment and decrement of an integer
- May interfere with inlining optimizations

## Non-constrained

- A few additional function calls inserted between the intercepted call, and the real method being invoked
  - These functions are relatively light, checking state stored in integers

***When the sanitizer is enabled, there is a (minimal) runtime cost***

**RealtimeSanitizer has no  
runtime cost without  
*-fsanitize=realtime***



# Performance constraints code conversion cost

Comparing and contrasting

# Converting a codebase

```
void process() noexcept  
[[clang::nonblocking]]  
{  
    foo();  
}
```



```
void foo() noexcept  
[[clang::nonblocking]]  
{  
    bar();  
};
```

```
noexcept  
blocking]]
```

'locking' function  
non-'nonblocking'

function

Upfront effort to  
convert to perf.  
constraints attributes

But it's *worth it*

**No engineering cost  
for new projects**



1. Run time vs compile time
2. False negatives and false positives
- 3. Cost**
4. Disabling each tool
5. Using each tool outside of LLVM

**CONTRAST**

1. Run time vs compile time
2. False negatives and false positives
3. Cost
4. **Disabling each tool**
5. Using each tool outside of LLVM

# CONTRAST

## RealtimeSanitizer

Disable with  
`__rtsan::ScopedDisabler`

## Performance constraints

Disable using call-site macro

# Disabling function effects warnings

```
void process() noexcept  
[[clang::nonblocking]]  
{  
    foo();  
}
```

# Disabling function effects warnings

```
void process() noexcept
[[clang::nonblocking]]
{
    NONBLOCKING_UNSAFE(foo());
}
```



```
#define NONBLOCKING_UNSAFE(...) \
    _Pragma("clang diagnostic push") \
    _Pragma("clang diagnostic ignored \\"-Wunknown-warning-option\\") \
    _Pragma("clang diagnostic ignored \\"-Wfunction-effects\\") \
    __VA_ARGS__ \
    _Pragma("clang diagnostic pop")
```

## Disabling RTSan

```
#include <sanitizer/rtsan_interface.h>
```

```
void lock_error_mutex(std::mutex& m)
{
    __rtsan::ScopedDisabler disabler{};
    m.lock();
}
```



```
void process()
    noexcept [[clang::nonblocking]]
{
    if (buffer_overflow)
        lock_error_mutex(m);
}
```

1. Run time vs compile time
2. False negatives and false positives
3. Cost
4. **Disabling each tool**
5. Using each tool outside of LLVM

# CONTRAST

1. Run time vs compile time
2. False negatives and false positives
3. Cost
4. Disabling each tool
5. **Using each tool outside of LLVM**

# CONTRAST



## RealtimeSanitizer

Can be used with other  
compilers (with a little hack)

## Performance constraints

Dependent on clang 20



# Standalone RTSan

Comparing and contrasting



**realtime-sanitizer/rtsan**

```
#define __SANITIZE_REALTIME  
#include "rtsan_standalone/rtsan_standalone.h"
```

## Standalone RTSan

1. #define \_\_SANITIZE\_REALTIME to conditionally enable the sanitizer
2. Find and include rtsan\_standalone.h

# Standalone RTSan

```
#define __SANITIZE_REALTIME
#include "rtsan_standalone/rtsan_standalone.h"
```

```
void main() {
    __rtsan::Initialize();
    ...
}

void process() {
    __rtsan::ScopedSanitizeRealtime ssr;

    auto x = std::make_unique<float>(3.0f);
}
```

1. #define \_\_SANITIZE\_REALTIME to conditionally enable the sanitizer
2. Find and include rtsan\_standalone.h
3. Initialize rtsan early in your process
4. Insert ScopedSanitizeRealtime in any blocks that would be [[nonblocking]]

# Standalone RTSan

```
#define __SANITIZE_REALTIME
#include "rtsan_standalone/rtsan_standalone.h"
```

```
void main() {
    __rtsan::Initialize();
    ...
}

void process() {
    __rtsan::ScopedSanitizeRealtime ssr;

    auto x = std::make_unique<float>(3.0f);
}
```

```
// CMakeLists.txt
```

```
target_link_libraries(helloWorld
    PRIVATE libclang_rt.rtsan_osx_dynamic.dylib
)
```

1. #define \_\_SANITIZE\_REALTIME to conditionally enable the sanitizer
2. Find and include rtsan\_standalone.h
3. Initialize rtsan early in your process
4. Insert ScopedSanitizeRealtime in any blocks that would be [[nonblocking]]
5. Build and link the rtsan runtime from LLVM

# Standalone RTSan

```
#define __SANITIZE_REALTIME
#include "rtsan_standalone/rtsan_standalone.h"
```

```
void main() {
    __rtsan::Initialize();
    ...
}

void process() {
    __rtsan::ScopedSanitizeRealtime ssr;

    auto x = std::make_unique<float>(3.0f);
}
```

```
// CMakeLists.txt
```

```
target_link_libraries(helloWorld
    PRIVATE libclang_rt.rtsan_osx_dynamic.dylib
)
```

1. #define \_\_SANITIZE\_REALTIME to conditionally enable the sanitizer
2. Find and include rtsan\_standalone.h
3. Initialize rtsan early in your process
4. Insert ScopedSanitizeRealtime in any blocks that would be [[nonblocking]]
5. Build and link the rtsan runtime from LLVM

## Standalone RTSan

AppleClang 15 

GNU 14 on Linux 



`realtime-sanitizer/rtsan`





# Summary



**How do I use performance  
constraints attributes?**

Summary

# Summary: performance constraints attributes

```
float process (float x) noexcept [[clang::nonblocking]]
{
    auto y = std::vector<float> (16);
    ...
}
```

```
> clang++ -Wfunction-effects -Wperf-constraint-implies-noexcept main.cpp
```

```
main.cpp:22:8: error: 'nonblocking' function must not call non-'nonblocking' function
'std::vector<float>::~~vector' [-Werror,-Wfunction-effects]
```

```
22 |     auto y = std::vector<float> (16);
    |           ^
```

```
main.cpp:22:12: note: in template expansion here
```

```
22 |     auto y = std::vector<float> (16);
    |           ^
```

```
main.cpp:22:12: error: 'nonblocking' function must not call non-'nonblocking' function
'std::vector<float>::vector' [-Werror,-Wfunction-effects]
```

```
22 |     auto y = std::vector<float> (16);
    |           ^
```



# How do I use RealtimeSanitizer?

Summary

# Summary: RealtimeSanitizer

```
float process (float x) noexcept [[clang::nonblocking]]
{
    auto y = std::vector<float> (16);
    ...
}
```

```
> clang++ -fsanitize=realtime main.cpp
```

```
> ./a.out
```

```
==86660==ERROR: RealtimeSanitizer: unsafe-library-call
```

```
Intercepted call to real-time unsafe function `malloc` in real-time context!
```

```
#0 0x103184cfc in malloc rtsan_interceptors.cpp:225
```

```
#1 0x18cb16524 in operator new(unsigned long)+0x1c
```

```
#2 0x102c03590 in std::__1::__libcxx_allocate
```

```
...
```

```
#11 0x102c02b38 in process(float)+0x28 /app/example.cpp:6:14
```

```
#12 0x102c02c00 in main+0x1c /app/example.cpp:12:5
```

## RealtimeSanitizer

- Detect issues at run time
- Possible false negatives
- Run time cost
- Disable using ScopedDisabler
- Can be used with other compilers

## Both

Detect real-time safety issues

Use  
[[clang::nonblocking]]

## Perf. constraints

- Detect issues at compile time
- Possible false positives
- Code conversion cost
- Disabled by turning off warning at call site
- Dependent on clang 20+

**Both tools were  
designed to  
complement each  
other**

Use both, and write  
better real-time  
code!



# Thanks!

- Doug Wyatt
- Ali Barker
- Reviews on LLVM work
  - @vitalybuka
  - @maskray
  - @sirride
  - @AaronBallman
  - ... many more ...
- Slide review and contributions
  - Ryan Avery
  - Oliver Graff
  - Stuart Glaser

Three thick diagonal lines are positioned around the central text. A cyan line runs from the top-left towards the bottom-left. A yellow line runs from the top-center towards the middle-right. A magenta line runs from the bottom-center towards the top-right, intersecting the yellow line.

**Questions?**



**Real-time programming meetup**

**CppCon 2024**

**6pm Wednesday**

**Mountain Pass Sports Bar**

**Anyone welcome!**

# Help improve RealtimeSanitizer

- Adding support
  - Windows
  - Android
  - Other architectures (than x86-64 and arm64)
- Adding new interceptors
- Testing out experimental unbound loop checking on your codebase
- General feedback/testing

We are happy to help you get started!  
Find us on Discord



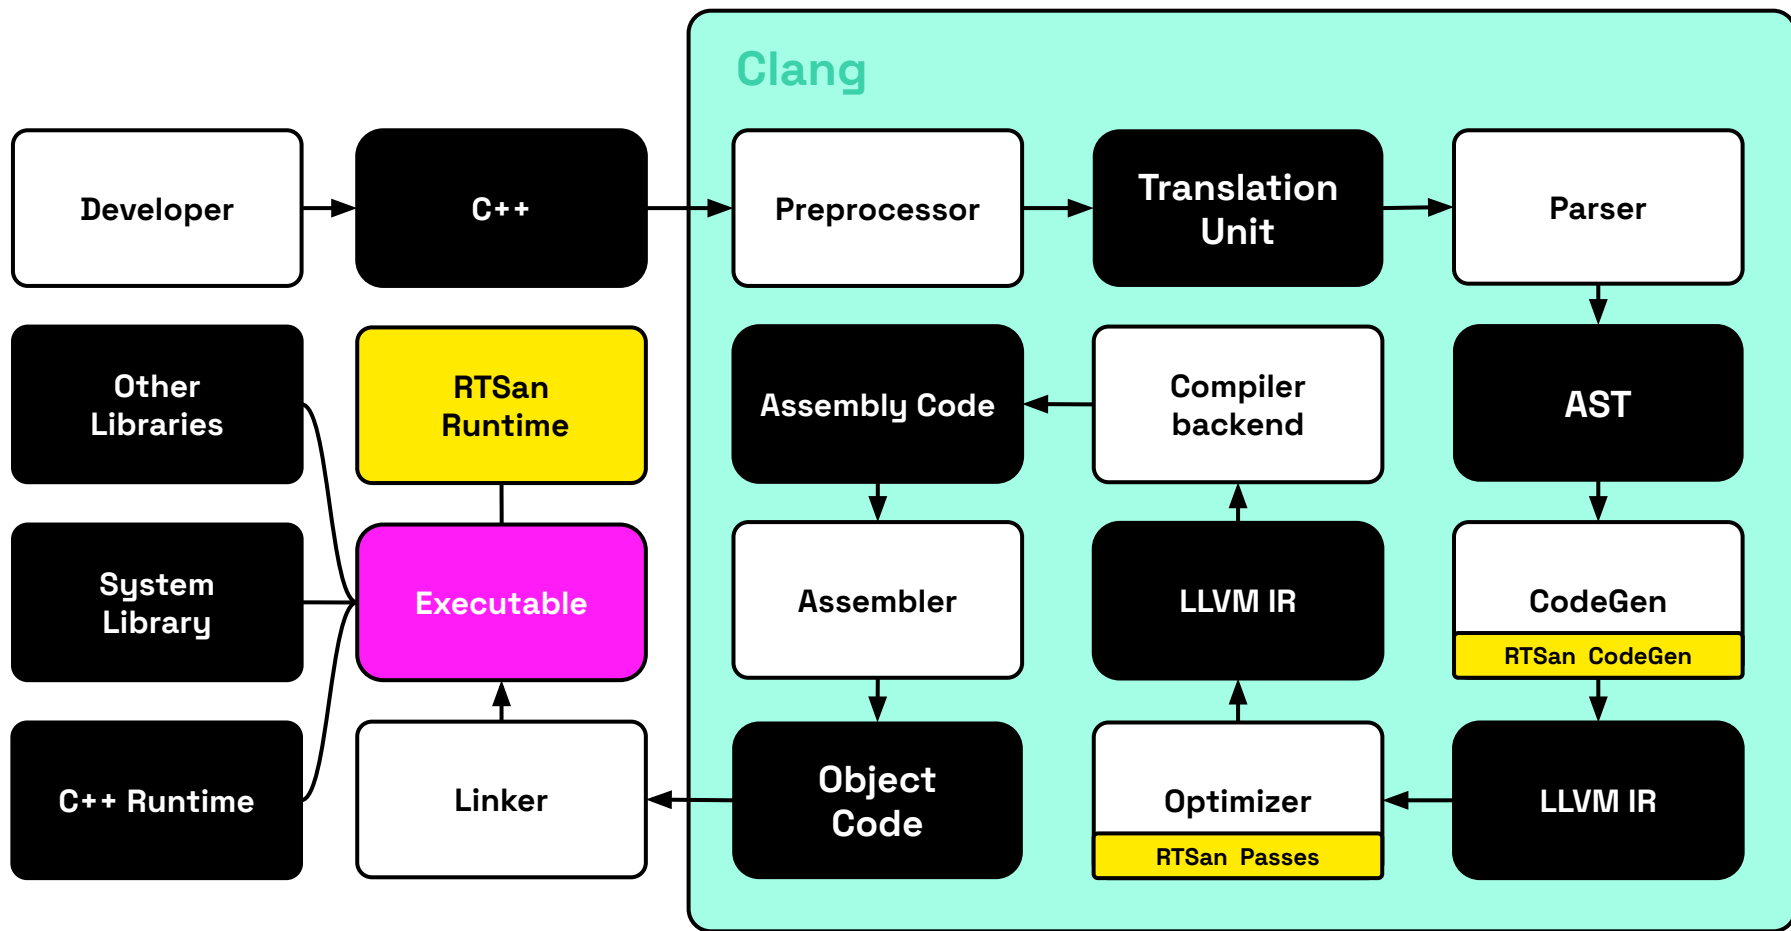
# Appendix

# What about DTrace?

- DTrace does a lot of similar things to RTSan
- However we wanted to push for RTSan in a central repository, where there was “one source of truth” for real-time safety
  - Duplication of “intercepted functions” across companies
  - Different norms marking what is real-time context
    - Can do this with function names, or thread priority
- DTrace also requires bypassing security on OSX, which is sometimes disallowed in some IT departments
- We want every platform to be supported trivially
- Sanitizer approach offers extra benefits like scoped disabling

# RealtimeSanitizer road map

- Soon
  - More interceptors
  - User defined `[[blocking]]` functions
  - Better interop with existing sanitizers (`RTSAN_OPTIONS`)
  - Adding “continue mode”
  - Suppression lists
- Later
  - Detection of nondeterministic loops
- Someday?
  - Windows support
  - Android/iOS support
  - Other architectures beyond Arm64 and x86\_64







**Customization**

The image features the word "Customization" in a bold, black, sans-serif font, centered horizontally. It is surrounded by three thick, solid-colored lines: a cyan line on the left, a yellow line on the top right, and a magenta line on the bottom right. The lines are oriented diagonally, creating a dynamic, abstract composition.


**“system call” ?**



# What is a system call?

RealtimeSanitizer

# A system call...



- Is a request to the operating system for interaction with a shared resource
- Has a specific number associated with them, which is passed to the processor via a register.
- Is not (typically) made directly from user code, but made from a wrapper library like libc.

# macOS

Interception is directly supported on macOS, using a feature called **dyld interposing**.

When loading any Mach-O binary, **dyld** treats:

- a pair of function pointers in a `__DATA, __interpose` section...
- ...as an instruction to **re-route all calls** from one function to the other

Calling the original function is simply a case of calling the original!

```
struct interpose_substitution
{
    const size_t ptr_to_replacement_function;
    const size_t ptr_to_original_function;
};

INTERCEPTOR (void *, malloc, size_t size)
{
    log_system_call ("malloc");
    return REAL (malloc) (size);
}

// expands to...

extern "C" void * malloc (size_t size);
extern "C" void * wrap_malloc (size_t size);

const interpose_substitution substitution_malloc[]
    __attribute__((section("__DATA, __interpose"))) = {
    {
        reinterpret_cast<const size_t> (wrap_malloc),
        reinterpret_cast<const size_t> (malloc),
    },
};

extern "C" void * wrap_malloc (size_t size)
{
    log_system_call ("malloc");
    return malloc (size);
}
```

# Linux

It's more straightforward to replace a function on Linux. If the function symbol is defined in your executable, that'll get found "first".

But calling the original (replaced) function is a little trickier!

It's possible to use **dlsym** to find the "next" occurrence of the function symbol in the search order.

Must be assigned before first call to malloc!



```
INTERCEPTOR (void *, malloc, size_t size)
{
    log_system_call ("malloc");
    return REAL (malloc) (size);
}
```

*// expands to...*

```
typedef void * (*malloc_function_type) (size_t size);

namespace __interception {
    malloc_function_type real_malloc;
}
```

```
void * malloc (size_t size)
{
    log_system_call ("malloc");
    return __interception::real_malloc (size);
}
```

```
INITIALISE_INTERCEPTOR (malloc);
```

*// expands to...*

```
__interception::real_malloc = dlsym (RTLD_NEXT, "malloc");
```

# C++



# LLVM IR



# Binary

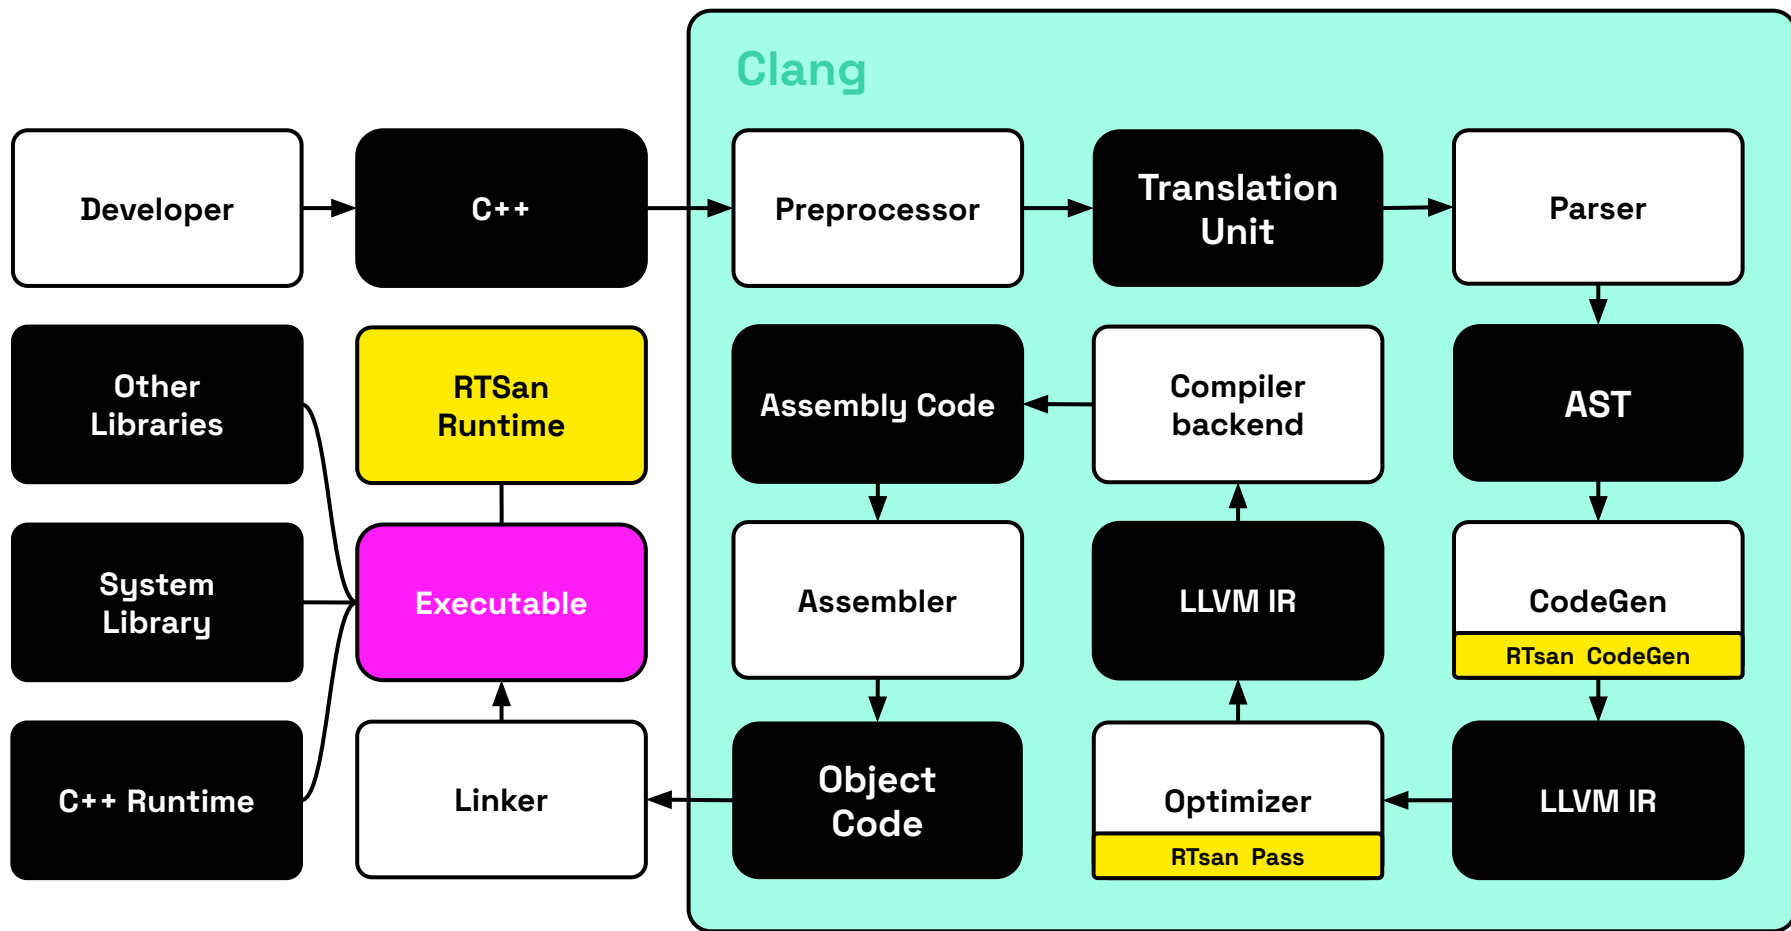
```
auto exit_code()
{
    return 0;
}

int main()
{
    return exit_code();
}
```

```
define i32 @_Z9exit_codev() #0 {
    ret i32 0
}

define i32 @main() #1 {
    %1 = alloca i32, align 4
    store i32 0, ptr %1, align 4
    %2 = call i32 @_Z9exit_codev()
    ret i32 %2
}
```

```
cffa edfe 0c00 0001 0000 0000 0200 0000 .....
1000 0000 e802 0000 8500 2000 0000 0000 .....
1900 0000 4800 0000 5f5f 5041 4745 5a45 .....H...PAGEZE
524f 0000 0000 0000 0000 0000 0000 0000 R0.....
0000 0000 0100 0000 0000 0000 0000 0000 .....
0000 0000 0000 0000 0000 0000 0000 0000 .....
0000 0000 0000 0000 1900 0000 e800 0000 .....
5f5f 5445 5854 0000 0000 0000 0000 0000 __TEXT.....
0000 0000 0100 0000 0040 0000 0000 0000 .....a.....
0000 0000 0000 0000 0040 0000 0000 0000 .....a.....
0500 0000 0500 0000 0200 0000 0000 0000 .....
5f5f 7465 7874 0000 0000 0000 0000 0000 __text.....
5f5f 5445 5854 0000 0000 0000 0000 0000 __TEXT.....
783f 0000 0100 0000 2800 0000 0000 0000 x?.....(.....
783f 0000 0200 0000 0000 0000 0000 0000 x?.....
0004 0080 0000 0000 0000 0000 0000 0000 __unwind_info...
5f5f 756e 7769 6e64 5f69 6e66 6f00 0000 __TEXT.....
5f5f 5445 5854 0000 0000 0000 0000 0000 .?.....
a03f 0000 0100 0000 0000 0000 0000 0000 .?.....
a03f 0000 0200 0000 0000 0000 0000 0000 .?.....
0000 0000 0000 0000 0000 0000 0000 0000 .....
1900 0000 4800 0000 5f5f 4c49 4e4b 4544 .....H...LINKED
4954 0000 0000 0000 0040 0000 0100 0000 IT.....a.....
0040 0000 0000 0000 0040 0000 0000 0000 .a.....a.....
0802 0000 0000 0000 0100 0000 0100 0000 .....4.....
0000 0000 0000 0000 3400 0080 1000 0000 .a..8...3.....
0040 0000 3800 0000 3300 0080 1000 0000 8a..H.....
3840 0000 4800 0000 0200 0000 1800 0000 .a.....a..0...
8840 0000 0300 0000 b340 0000 3000 0000 .....P.....
0b00 0000 5000 0000 0000 0000 0000 0000 .....
0000 0000 0300 0000 0300 0000 0000 0000 .....
0000 0000 0000 0000 0000 0000 0000 0000 .....
0000 0000 0000 0000 0000 0000 0000 0000 .....
0000 0000 0000 0000 0000 0000 0000 0000 .....
0e00 0000 2000 0000 0c00 0000 2f75 7372 ....../usr
2f6c 6962 2f64 796c 6400 0000 0000 0000 /lib/dyld.....
1b00 0000 1800 0000 80ae 897c 33d6 3c11 .....|3.<..
843c a99f d7b9 9f38 3200 0000 2000 0000 .<.....82...
0100 0000 0000 0e00 0000 0e00 0100 0000 .....
0300 0000 0007 f703 2a00 0000 1000 0000 .....*.....
0000 0000 0000 0000 2800 0080 1800 0000 .....(.....
803f 0000 0000 0000 0000 0000 0000 0000 .?.....
```





**By detecting calls to  
libc wrapper functions,  
we can warn about  
potential syscalls**